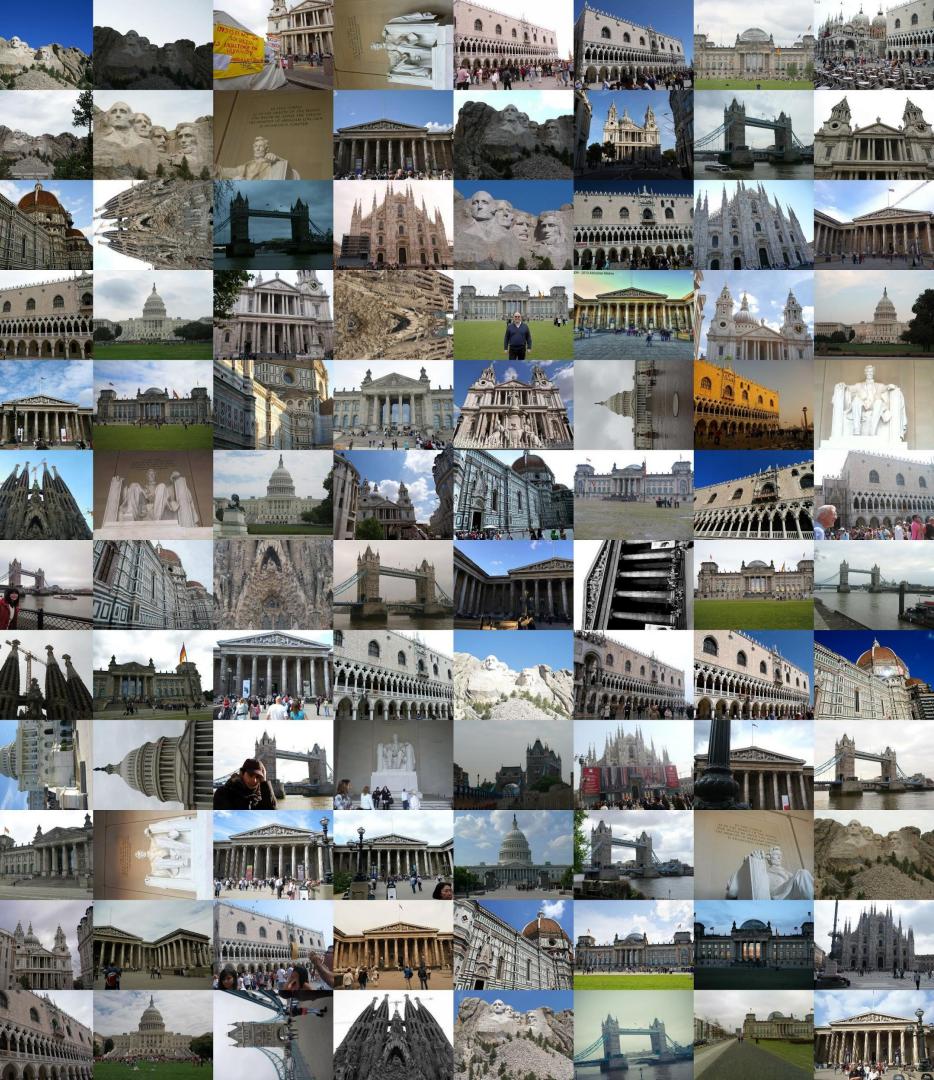


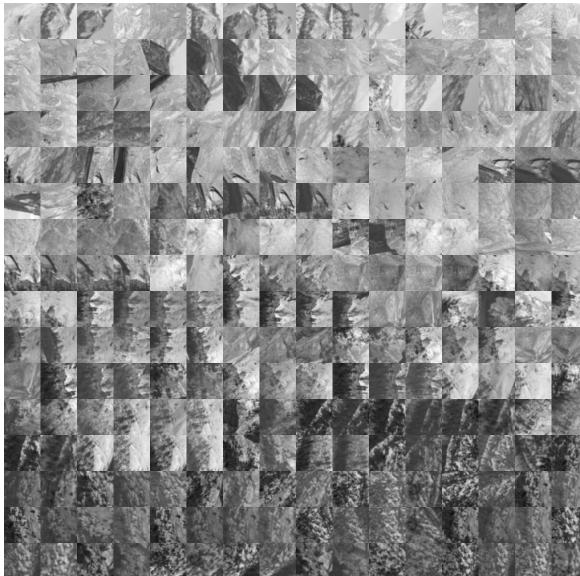
Phototourism Challenge

Eduard Trulls (Google)
Kwang Moo Yi (U. Victoria)
Sri Raghu Malireddi (U. Victoria)
Yuhe Jin (U. Victoria)

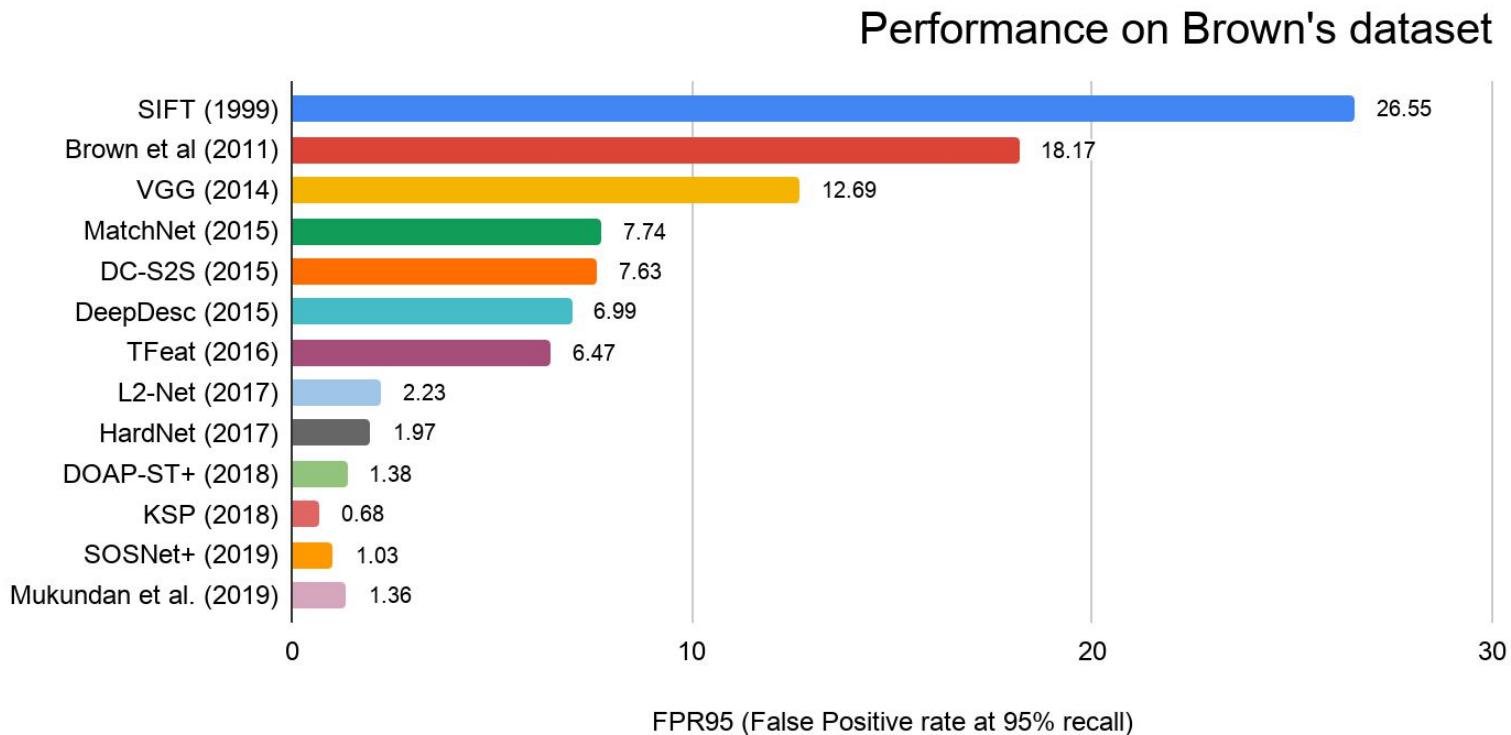


How good is
<insert-your-favorite-method-here>
in practice?

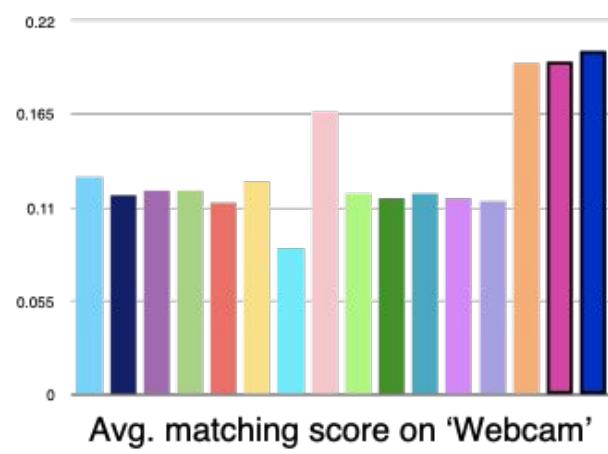
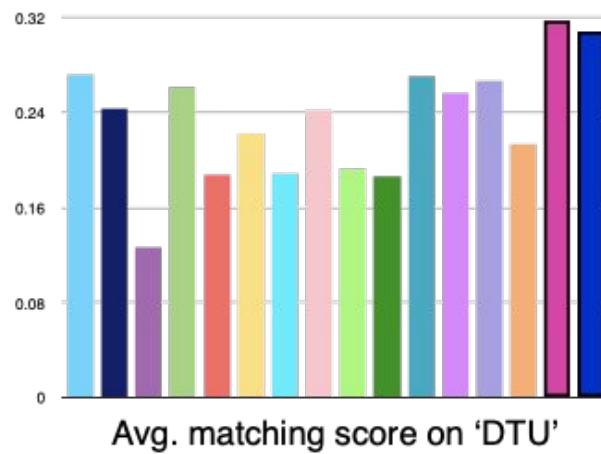
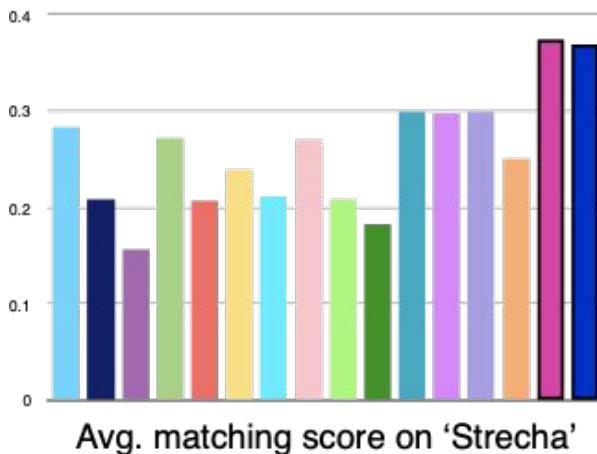
Current benchmarks are **saturated**



Current benchmarks are saturated

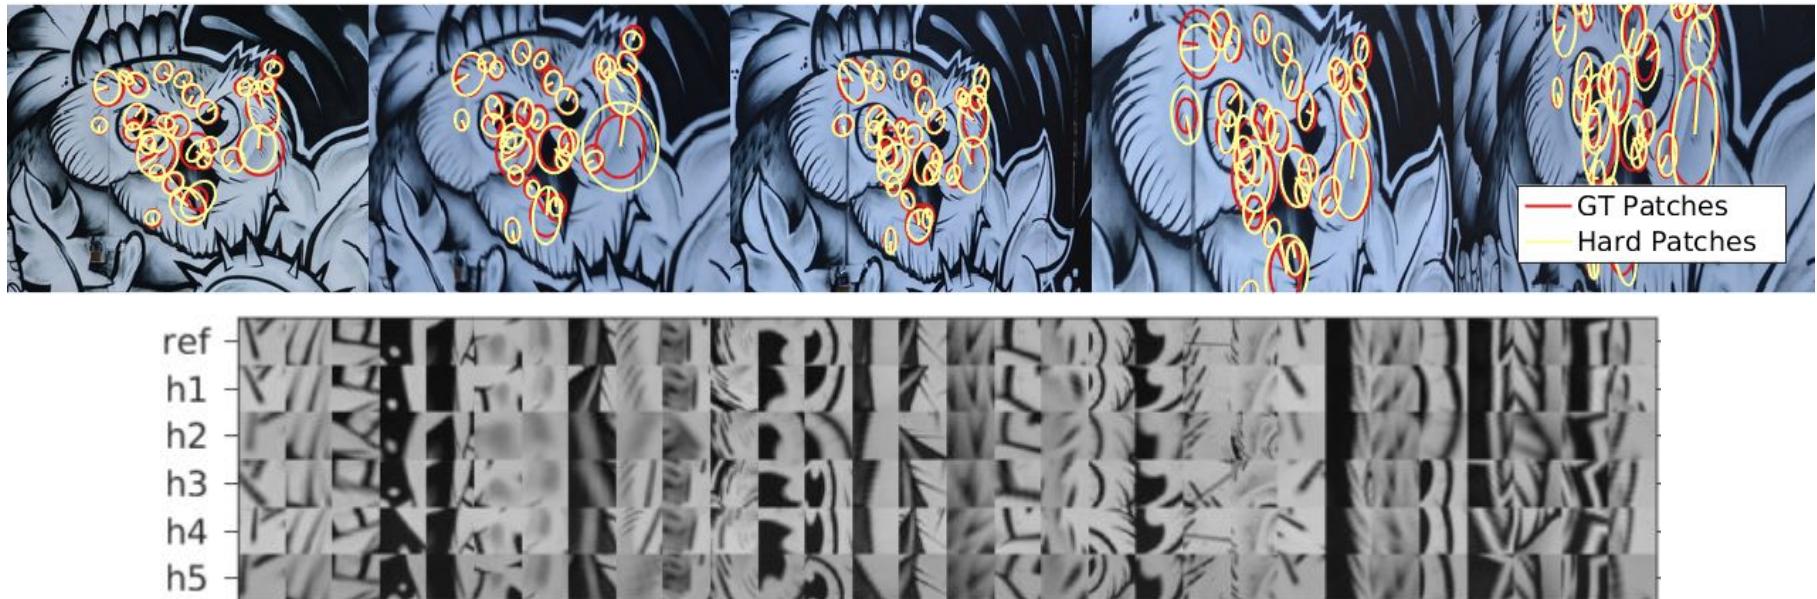


Current benchmarks are **not** representative



Towards proper benchmarking -- H(omography)Patches

Task: patch matching under affine transformation or illumination changes



[HPatches: A benchmark and evaluation of handcrafted and learned local descriptors](#). V. Balntas et al., CVPR'17

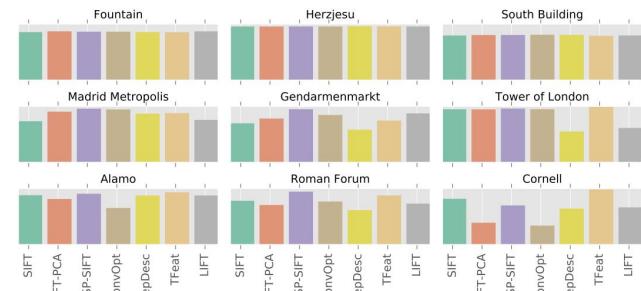
Source: github.com/hpatches/hpatches-dataset

Towards proper benchmarking – SfM (COLMAP)

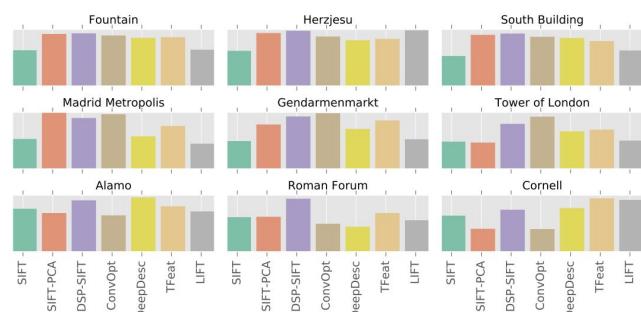
Task: 3D reconstruction with local features



Number of registered images



Number of registered 3D points



Comparative Evaluation of Hand-Crafted and Learned Local Features.

Schönberger et al., CVPR'17. Source: github.com/ahojnnes/local-feature-evaluation

Depth comes at a cost



[On benchmarking camera calibration and multi-view stereo for high resolution imagery](#). Strecha et al., CVPR'08.

How good is
<insert-your-favorite-method-here>
in practice?

How good is

<insert-your-favorite-method-here>

in practice?

Towards practical evaluation

- Variation + Volume

Towards practical evaluation

- Variation + Volume
 - Phototourism data: viewpoint, sensors, illumination, motion blur, occlusions, etc
 - Large-scale: ~30k images
 - Images, poses & depth: suitable for multiple tasks

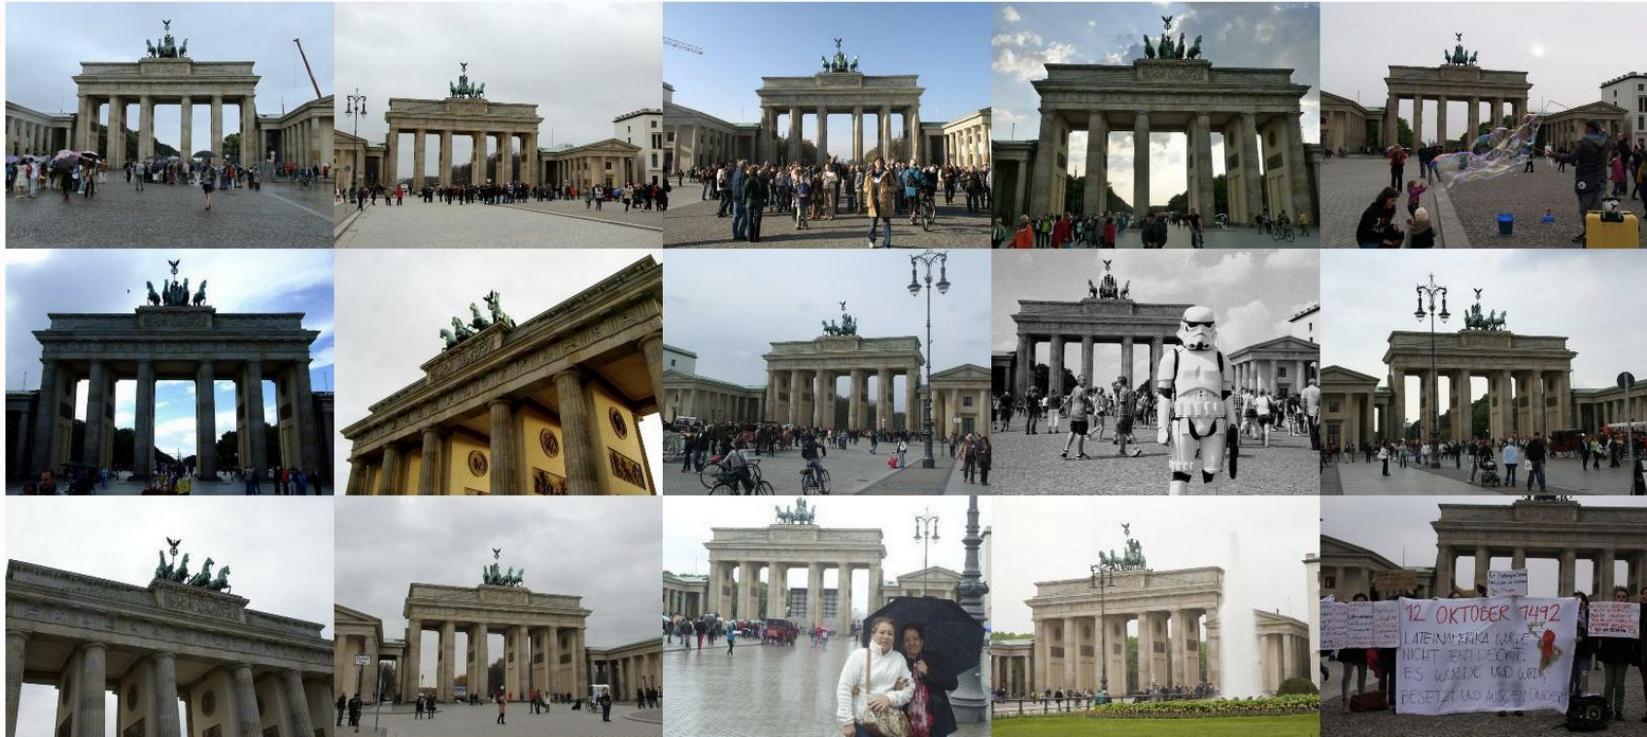
Towards practical evaluation

- Variation + Volume
 - Phototourism data: viewpoint, sensors, illumination, motion blur, occlusions, etc
 - Large-scale: ~30k images
 - Images, poses & depth: suitable for multiple tasks
- Image-level evaluation
 - Matching scores

Towards practical evaluation

- Variation + Volume
 - Phototourism data: viewpoint, sensors, illumination, motion blur, occlusions, etc
 - Large-scale: ~30k images
 - Images, poses & depth: suitable for multiple tasks
- Image-level evaluation
 - Matching scores
 - Stereo: Camera pose accuracy
 - SfM: Camera pose accuracy + Metrics by Schönberger et al. CVPR'17

The phototourism challenge: Data



Examples from brandenburg_gate

The phototourism challenge: Data



The phototourism challenge: Data

Training sequences	▲	Num. images	◆	Num. 3D points	◆
brandenburg_gate		1363		100040	
buckingham_palace		1676		234052	
colosseum_exterior		2063		259807	
grand_place_brussels		1083		229788	
hagia_sophia_interior		888		235541	
notre_dame_front_facade		3765		488895	
palace_of_westminster		983		115868	
pantheon_exterior		1401		166923	
prague_old_town_square		2316		558600	
sacre_coeur		1179		140659	
st_peters_square		2504		232329	
taj_mahal		1312		94121	
temple_nara_japan		904		92131	
trevi_fountain		3191		580673	
westminster_abbey		1061		198222	
Total		25.6k		3.7M	

- 25k images in total for training.
- “Quasi” ground truth data is generated by performing SfM with COLMAP with all images.
 - Assumption: Images registered in COLMAP are accurate given enough images.
- Valid pairs are generated via simple visibility check.

The phototourism challenge: Data

Test sequences	▲	Num. images	◆	Num. 3D points	◆
british_museum		660		73569	
florence_cathedral_side		108		44143	
lincoln_memorial_statue		850		58661	
london_bridge		629		72235	
milan_cathedral		124		33905	
mount_rushmore		138		45350	
piazza_san_marco		249		95895	
reichstag		75		17823	
sagrada_familia		401		120723	
st_pauls_cathedral		615		98872	
united_states_capitol		258		35095	
Total		4107		696k	

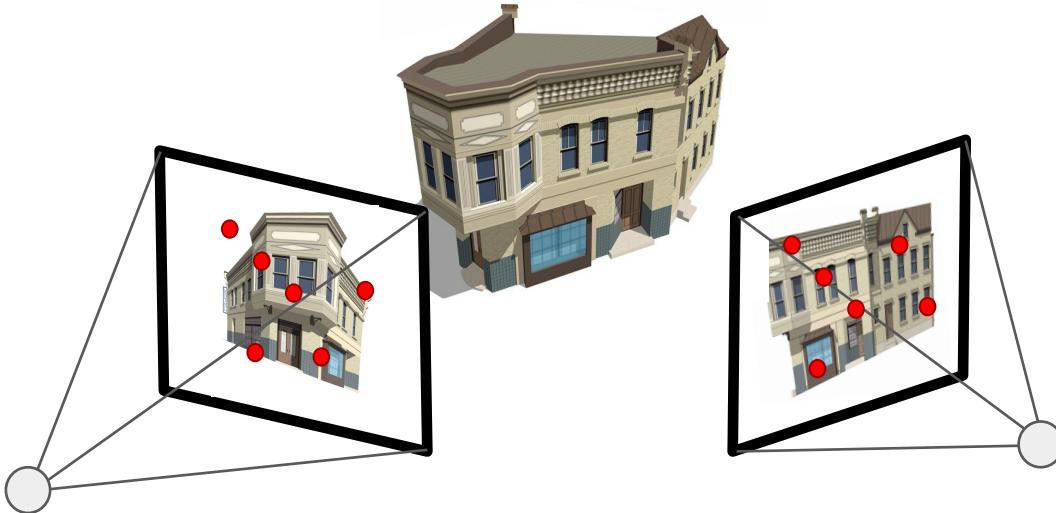
- 4k images in total for testing.
- Random bags of images are subsampled to form test subsets (size: 3, 5, 10, 25).

The phototourism challenge: local features



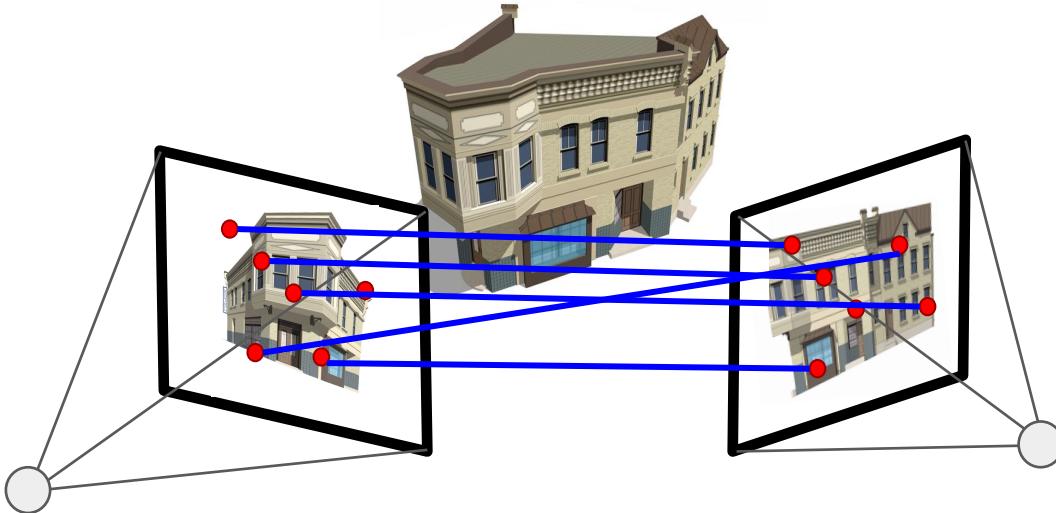
- Submission: **Features**
- IMW evaluates them via a typical stereo/SfM pipeline
 - Nearest neighbor matching
 - 1-to-1 matching
 - RANSAC
 - COLMAP

The phototourism challenge: local features



- Submission: **Features**
- IMW evaluates them via a typical stereo/SfM pipeline
 - Nearest neighbor matching
 - 1-to-1 matching
 - RANSAC_F
 - COLMAP

The phototourism challenge: matches



- Submission: **Features + Matches**
- IMW evaluates them via a typical stereo/SfM pipeline
 - ~~Nearest neighbor matching~~
 - ~~1 to 1 matching~~
 - RANSAC_F
 - COLMAP

The phototourism challenge: poses



- Submission: **Poses**
- IMW evaluates them via a typical stereo/SfM pipeline
 - ~~Nearest neighbor matching~~
 - ~~1 to 1 matching~~
 - ~~RANSAC_F~~
 - ~~COLMAP~~

The phototourism challenge: Stereo

Stereo – averaged over all sequences									
Method	Date	Type	#kp	MS	mAP 5°	mAP 10°	mAP 15°	mAP 20°	mAP 25°
+SIFT (OpenCV) kp:8000, match:nn	19-04-24	F	7884.4	0.2148	0.0004	0.0068	0.0277	0.0692	0.1303
+SIFT (OpenCV) kp:2048, match:nn	19-05-17	F	2048.0	0.2215	0.0004	0.0055	0.0223	0.0570	0.1060

Matching score, but with
symmetric epipolar
distance for thresholding.

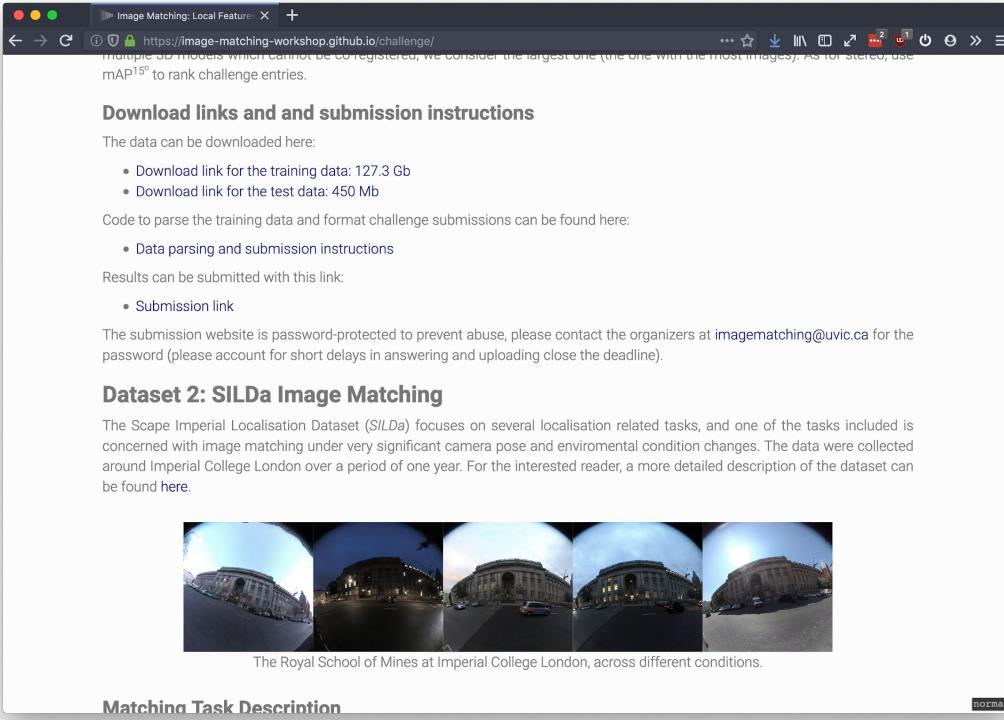
Mean average precision -- average
ratio of correct estimates under
varying thresholds until 15 degrees
(considering both R, t)

The phototourism challenge: SfM

MVS – averaged over all sequences												
Method	Date	Type	Ims (%)	#Pts	SR	TL	mAP ^{5°}	mAP ^{10°}	mAP ^{15°}	mAP ^{20°}	mAP ^{25°}	ATE
SIFT (OpenCV) kp:8000, match:nn	19-04-24	F	93.6	4341.5	88.5	3.15	0.2881	0.3640	0.4146	0.4550	0.4901	—
SIFT (OpenCV) kp:2048, match:nn	19-05-17	F	85.3	1214.3	76.9	2.93	0.1521	0.2060	0.2439	0.2762	0.3038	—

Mean average precision -- average ratio of correct estimates under varying thresholds until 15 degrees (considering both R, t)

The phototourism challenge: Submission



The screenshot shows a web browser window titled "Image Matching: Local Feature". The URL is <https://image-matching-workshop.github.io/challenge/>. The page content includes:

- Download links and submission instructions**:
The data can be downloaded here:
 - Download link for the training data: 127.3 Gb
 - Download link for the test data: 450 Mb
- Code to parse the training data and format challenge submissions can be found here:
 - Data parsing and submission instructions
- Results can be submitted with this link:
 - Submission link
- The submission website is password-protected to prevent abuse, please contact the organizers at imagematching@uvic.ca for the password (please account for short delays in answering and uploading close the deadline).

Dataset 2: SILDa Image Matching

The Scape Imperial Localisation Dataset (*SILDa*) focuses on several localisation related tasks, and one of the tasks included is concerned with image matching under very significant camera pose and environmental condition changes. The data were collected around Imperial College London over a period of one year. For the interested reader, a more detailed description of the dataset can be found [here](#).



The Royal School of Mines at Imperial College London, across different conditions.

Matching Task Description

- Upload server is password protected
 - Contact us for password
- Submission rules to be updated soon
 - We used roughly 55 core-years for this year challenge alone :-)



- Code release soon
 - Welcoming contributions (and criticism!)

SILDa Challenge

Vassileios Balntas (Scape)



SILDa Challenge

Vassileios Balntas (Scape)

Axel Baroso (Imperial College London)
Krystian Mikolajczyk (Imperial College London)
Rigas Kouskouridas (Scape Technologies)
Duncan Frost (Scape Technologies)
Huub Heijnen (Scape Technologies)

SILDa: Key facts

- 14k images collected around Imperial College London over 1.5 year
- Rain, snow, sun, evening, night, morning
- Significant variations in the scenes







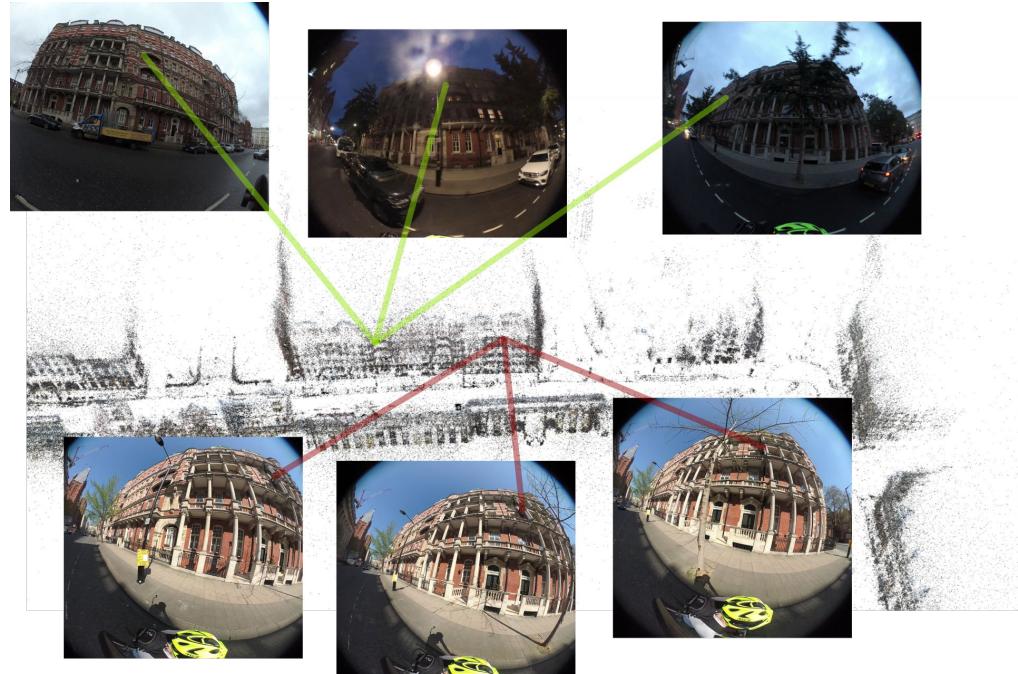
3D Reconstruction

- SfM with calibrated spherical cameras
- Chain SfM to help out matches: e.g. day -> evening & evening -> night.
- 1.4M points in the point cloud
- Covering almost 20 passes of 1.6km road



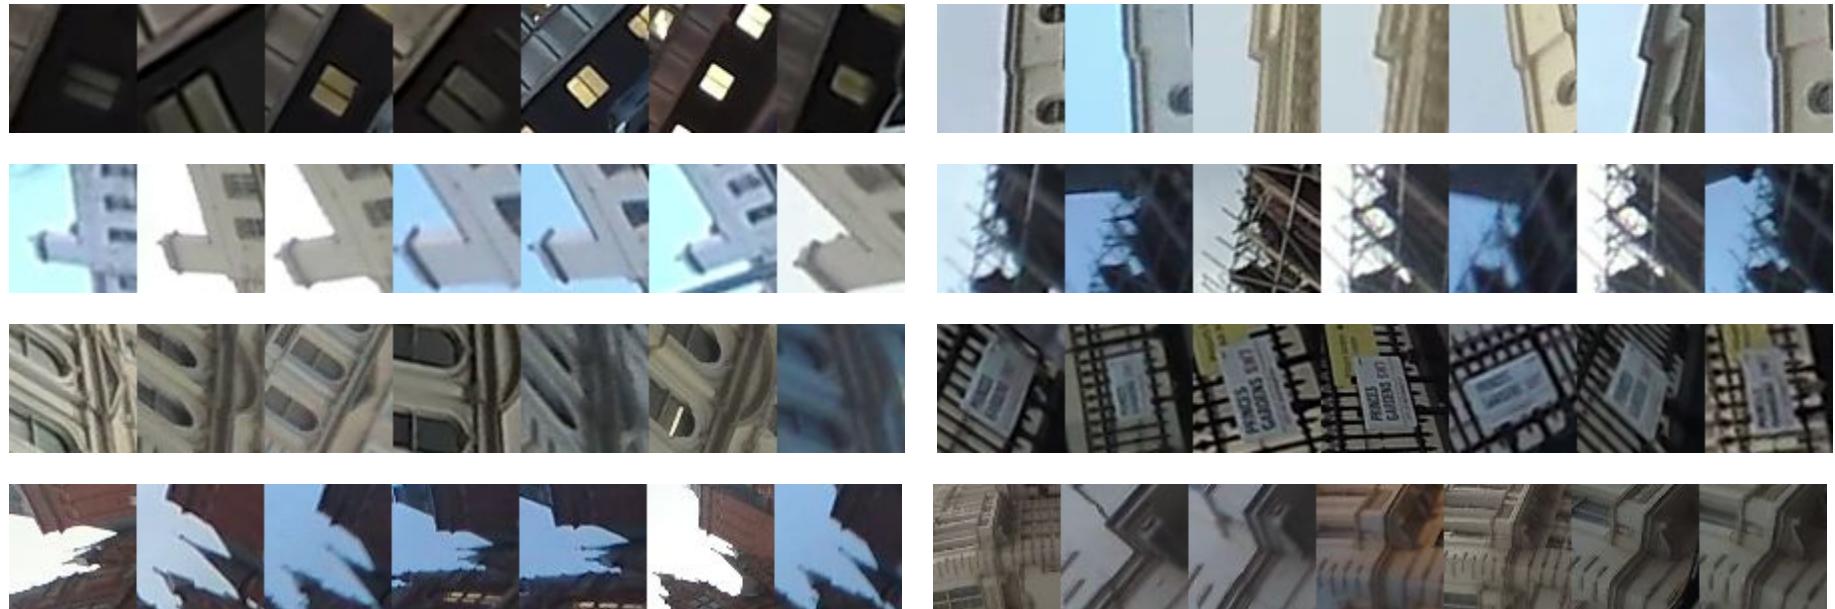
Local patches

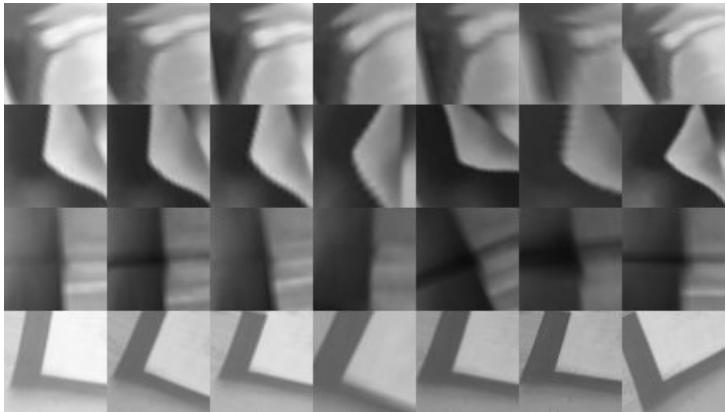
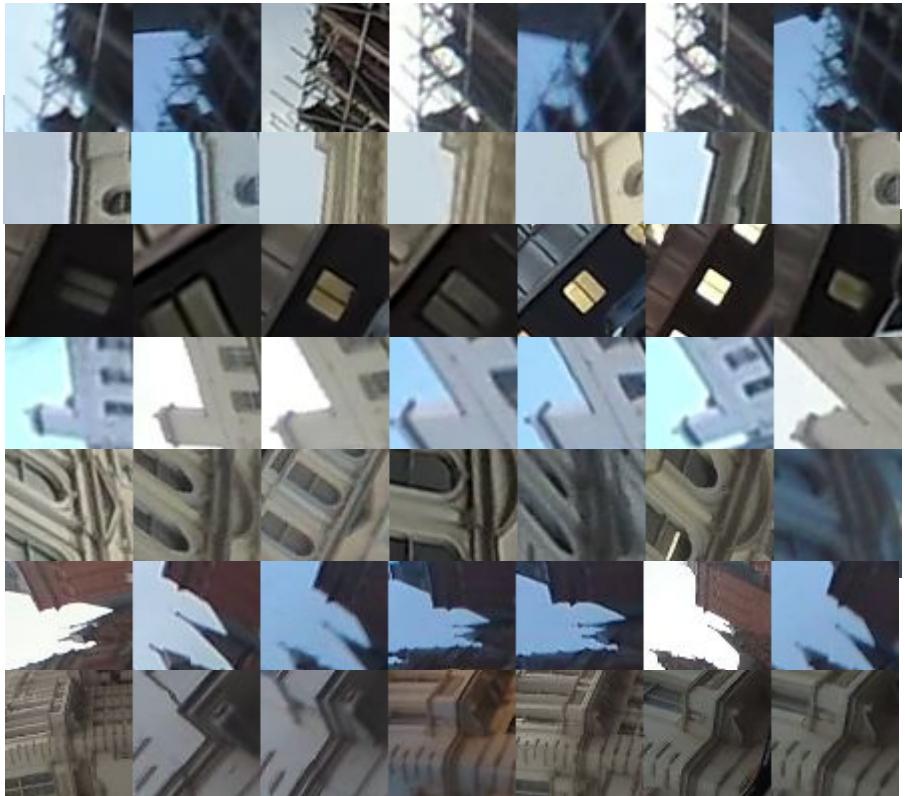
- Similarly to Brown and HPatches we extract a set of patches from the 3d points across different days, times and conditions



Local patches

- Similarly to Brown and HPatches we extract a set of patches from the 3d points across different days, times and conditions





Are patches still relevant?

- Is colour important for descriptors (CNN)?
- Is patch matching a good proxy for image matching?
- Is the separate evaluation of detector/descriptor the best strategy?

IMW Challenge: Image Pairs

- We generate 100k image pairs, which are deemed *difficult*
 - *difficult*: small number of inlier matches (<100) during the SfM process, but contain common point cloud points.
 - why focus on *difficult*?
 - *lots of SfM pairs are very incremental in terms of camera motion and end up having a big amount of inliers (>1000)*









Evaluation Protocol: Epipolar Arcs



Evaluation Protocol: Epipolar Arcs



SILDa challenge: Submission

Image Matching: Local Feature x +

https://image-matching-workshop.github.io/challenge/

Multiple SILDa entries will be accepted. We consider the largest one (the one with the most images). As for stereo, use mAP^{15°} to rank challenge entries.

Download links and submission instructions

The data can be downloaded here:

- Download link for the training data: 127.3 Gb
- Download link for the test data: 450 Mb

Code to parse the training data and format challenge submissions can be found here:

- Data parsing and submission instructions

Results can be submitted with this link:

- Submission link

The submission website is password-protected to prevent abuse, please contact the organizers at imagematching@uvic.ca for the password (please account for short delays in answering and uploading close the deadline).

Dataset 2: SILDa Image Matching

The Scape Imperial Localisation Dataset (*SILDa*) focuses on several localisation related tasks, and one of the tasks included is concerned with image matching under very significant camera pose and environmental condition changes. The data were collected around Imperial College London over a period of one year. For the interested reader, a more detailed description of the dataset can be found [here](#).



The Royal School of Mines at Imperial College London, across different conditions.

Matching Task Description

normal

- Online server will be available later on
- Hidden test set
- Future: more baselines D2Net, ContextDesc etc...

SILDa Matching Challenge: 3 Evaluation Metrics

- **Matching Scores:** Define a threshold on epipolar arc distance error, and use this to compute correct matches
- **Epipolar Arc Distance Statistics:** average/median epipolar arc distances between matches
- **Number of image pairs with more than 8 inliers**

Program

8:45 - 9:00	Welcome
9:00 - 9:30	Amir Zamir (Stanford/UC Berkeley) <i>Collection of Large-scale Densely-labeled 3D Data from the Real World Without a Single Click</i>
9:30 - 10:15	Jiri Matas (CTU Prague) <i>On the Art of Establishing Correspondence</i>
10:15 - 11:00	Coffee Break + Poster Session
11:15 - 12:00	Torsten Sattler (Chalmers U. of Technology, Gothenburg) <i>In Defense of Local Features for Visual Localization</i>
12:00 - 12:15	IMW2019 Challenge
12:15 - 12:30	Zixin Luo (HKUST) Winner of the Phototourism Challenge
12:30 - 12:45	Challenge results and awards